

# sklearn.feature\_extraction.text.TfidfVectorizer¶

 [scikit-learn.org/stable/modules/generated/sklearn.feature\\_extraction.text.TfidfVectorizer.html](https://scikit-learn.org/stable/modules/generated/sklearn.feature_extraction.text.TfidfVectorizer.html)

Convert a collection of raw documents to a matrix of TF-IDF features.

Equivalent to CountVectorizer followed by TfidfTransformer.

Read more in the [User Guide](#).

**input** : string {'filename', 'file', 'content'}

*If 'filename', the sequence passed as an argument to fit is expected to be a list of filenames that need reading to fetch the raw content to analyze.*

*If 'file', the sequence items must have a 'read' method (file-like object) that is called to fetch the bytes in memory.*

*Otherwise the input is expected to be the sequence strings or bytes items are expected to be analyzed directly.*

**encoding** : string, 'utf-8' by default.

*If bytes or files are given to analyze, this encoding is used to decode.*

**decode\_error** : {'strict', 'ignore', 'replace'}

*Instruction on what to do if a byte sequence is given to analyze that contains characters not of the given encoding. By default, it is 'strict', meaning that a UnicodeDecodeError will be raised. Other values are 'ignore' and 'replace'.*

**strip\_accents** : {'ascii', 'unicode', None}

*Remove accents during the preprocessing step. 'ascii' is a fast method that only works on characters that have an direct ASCII mapping. 'unicode' is a slightly slower method that works on any characters. None (default) does nothing.*

**analyzer** : string, {'word', 'char'} or callable

*Whether the feature should be made of word or character n-grams.*

*If a callable is passed it is used to extract the sequence of features out of the raw, unprocessed input.*

**preprocessor** : callable or None (default)

Override the preprocessing (string transformation) stage while preserving the tokenizing and n-grams generation steps.

**tokenizer** : callable or None (default)

Override the string tokenization step while preserving the preprocessing and n-grams generation steps. Only applies if `analyzer == 'word'`.

**ngram\_range** : tuple (min\_n, max\_n)

The lower and upper boundary of the range of n-values for different n-grams to be extracted. All values of n such that  $\text{min\_n} \leq n \leq \text{max\_n}$  will be used.

**stop\_words** : string {'english'}, list, or None (default)

If a string, it is passed to `_check_stop_list` and the appropriate stop list is returned. 'english' is currently the only supported string value.

If a list, that list is assumed to contain stop words, all of which will be removed from the resulting tokens. Only applies if `analyzer == 'word'`.

If None, no stop words will be used. `max_df` can be set to a value in the range [0.7, 1.0) to automatically detect and filter stop words based on intra corpus document frequency of terms.

**lowercase** : boolean, default True

Convert all characters to lowercase before tokenizing.

**token\_pattern** : string

Regular expression denoting what constitutes a "token", only used if `analyzer == 'word'`. The default regexp selects tokens of 2 or more alphanumeric characters (punctuation is completely ignored and always treated as a token separator).

**max\_df** : float in range [0.0, 1.0] or int, default=1.0

When building the vocabulary ignore terms that have a document frequency strictly higher than the given threshold (corpus-specific stop words). If float, the parameter represents a proportion of documents, integer absolute counts. This parameter is ignored if vocabulary is not None.

**min\_df** : float in range [0.0, 1.0] or int, default=1

*When building the vocabulary ignore terms that have a document frequency strictly lower than the given threshold. This value is also called cut-off in the literature. If float, the parameter represents a proportion of documents, integer absolute counts. This parameter is ignored if vocabulary is not None.*

**max\_features** : int or None, default=None

*If not None, build a vocabulary that only consider the top max\_features ordered by term frequency across the corpus.*

*This parameter is ignored if vocabulary is not None.*

**vocabulary** : Mapping or iterable, optional

*Either a Mapping (e.g., a dict) where keys are terms and values are indices in the feature matrix, or an iterable over terms. If not given, a vocabulary is determined from the input documents.*

**binary** : boolean, default=False

*If True, all non-zero term counts are set to 1. This does not mean outputs will have only 0/1 values, only that the tf term in tf-idf is binary. (Set idf and normalization to False to get 0/1 outputs.)*

**dtype** : type, optional

*Type of the matrix returned by fit\_transform() or transform().*

**norm** : 'l1', 'l2' or None, optional

*Norm used to normalize term vectors. None for no normalization.*

**use\_idf** : boolean, default=True

*Enable inverse-document-frequency reweighting.*

**smooth\_idf** : boolean, default=True

*Smooth idf weights by adding one to document frequencies, as if an extra document was seen containing every term in the collection exactly once. Prevents zero divisions.*

**sublinear\_tf** : boolean, default=False

| *Apply sublinear tf scaling, i.e. replace tf with  $1 + \log(tf)$ .*

## Parameters:

---

**vocabulary\_** : dict

| *A mapping of terms to feature indices.*

**idf\_** : array, shape = [n\_features], or None

| *The learned idf vector (global term weights) when `use_idf` is set to True, None otherwise.*

**stop\_words\_** : set

| *Terms that were ignored because they either:*

- *occurred in too many documents (`max_df`)*
- *occurred in too few documents (`min_df`)*
- *were cut off by feature selection (`max_features`).*

| *This is only available if no vocabulary was given.*

## Attributes:

---

See also

`CountVectorizer`

Tokenize the documents and count the occurrences of token and return them as a sparse matrix

`TfidfTransformer`

Apply Term Frequency Inverse Document Frequency normalization to a sparse matrix of occurrence counts.

## Notes

The `stop_words_` attribute can get large and increase the model size when pickling. This attribute is provided only for introspection and can be safely removed using `delattr` or set to None before pickling.

## Methods

<code>build_analyzer()</code>	Return a callable that handles preprocessing and tokenization
<code>build_preprocessor()</code>	Return a function to preprocess the text before tokenization

<code>build_tokenizer()</code>	Return a function that splits a string into a sequence of tokens
<code>decode(doc)</code>	Decode the input into a string of unicode symbols
<code>fit</code> ( <code>raw_documents[, y]</code> )	Learn vocabulary and idf from training set.
<code>fit_transform</code> ( <code>raw_documents[, y]</code> )	Learn vocabulary and idf, return term-document matrix.
<code>get_feature_names()</code>	Array mapping from feature integer indices to feature name
<code>get_params([deep])</code>	Get parameters for this estimator.
<code>get_stop_words()</code>	Build or fetch the effective stop words list
<code>inverse_transform</code> ( <code>X</code> )	Return terms per document with nonzero entries in X.
<code>set_params</code> ( <code>\**params</code> )	Set the parameters of this estimator.
<code>transform</code> ( <code>raw_documents[, copy]</code> )	Transform documents to document-term matrix.

`__init__(input=u'content', encoding=u'utf-8', decode_error=u'strict', strip_accents=None, lowercase=True, preprocessor=None, tokenizer=None, analyzer=u'word', stop_words=None, token_pattern=u'(?u)\b\w+\b', ngram_range=(1, 1), max_df=1.0, min_df=1, max_features=None, vocabulary=None, binary=False, dtype=, norm=u'l2', use_idf=True, smooth_idf=True, sublinear_tf=False)`[\[source\]](#)[¶](#)

`build_analyzer()`[\[source\]](#)[¶](#)

Return a callable that handles preprocessing and tokenization

`build_preprocessor()`[\[source\]](#)[¶](#)

Return a function to preprocess the text before tokenization

`build_tokenizer()`[\[source\]](#)[¶](#)

Return a function that splits a string into a sequence of tokens

`decode(doc)`[\[source\]](#)[¶](#)

Decode the input into a string of unicode symbols

The decoding strategy depends on the vectorizer parameters.

`fit(raw_documents, y=None)`[\[source\]](#)[¶](#)

Learn vocabulary and idf from training set.

**raw\_documents** : iterable

| *an iterable which yields either str, unicode or file objects*

**Parameters:**

---

**Returns:**     **self** : TfidfVectorizer

---

`fit_transform(raw_documents, y=None)`[\[source\]](#)<sup>¶</sup>

Learn vocabulary and idf, return term-document matrix.

This is equivalent to fit followed by transform, but more efficiently implemented.

**raw\_documents** : iterable

| *an iterable which yields either str, unicode or file objects*

**Parameters:**

---

**X** : sparse matrix, [n\_samples, n\_features]

| *Tf-idf-weighted document-term matrix.*

**Returns:**

---

`get_feature_names()`[\[source\]](#)<sup>¶</sup>

Array mapping from feature integer indices to feature name

`get_params(deep=True)`[\[source\]](#)<sup>¶</sup>

Get parameters for this estimator.

**deep** : boolean, optional

| *If True, will return the parameters for this estimator and contained subobjects that are estimators.*

**Parameters:**

---

**params** : mapping of string to any

| *Parameter names mapped to their values.*

**Returns:**

---

`get_stop_words()`[\[source\]](#)<sup>¶</sup>

Build or fetch the effective stop words list

`inverse_transform(X)[source]`

Return terms per document with nonzero entries in X.

**Parameters:** **X** : {array, sparse matrix}, shape = [n\_samples, n\_features]

---

**X\_inv** : list of arrays, len = n\_samples

| *List of arrays of terms.*

**Returns:**

---

`set_params(**params)[source]`

Set the parameters of this estimator.

The method works on simple estimators as well as on nested objects (such as pipelines). The latter have parameters of the form `<component>__<parameter>` so that it's possible to update each component of a nested object.

**Returns:** **self** :

---

`transform(raw_documents, copy=True)[source]`

Transform documents to document-term matrix.

Uses the vocabulary and document frequencies (df) learned by fit (or fit\_transform).

**raw\_documents** : iterable

| *an iterable which yields either str, unicode or file objects*

**copy** : boolean, default True

| *Whether to copy X and operate on the copy or perform in-place operations.*

**Parameters:**

---

**X** : sparse matrix, [n\_samples, n\_features]

| *Tf-idf-weighted document-term matrix.*

**Returns:**

---